



Logistic regression with missing covariates—Parameter estimation, model selection and prediction within a joint-modeling framework

Wei Jiang^{a,*}, Julie Josse^a, Marc Lavielle^a, TraumaBase Group^b

^a Inria XPOP and CMAP, École Polytechnique, France

^b Hôpital Beaujon, APHP, France

ARTICLE INFO

Article history:

Received 19 February 2019

Received in revised form 29 November 2019

Accepted 21 December 2019

Available online 30 December 2019

Keywords:

Incomplete data

Observed likelihood

Metropolis–Hastings

Public health

ABSTRACT

Logistic regression is a common classification method in supervised learning. Surprisingly, there are very few solutions for performing logistic regression with missing values in the covariates. A complete approach based on a stochastic approximation version of the EM algorithm is proposed in order to perform statistical inference with missing values, including the estimation of the parameters and their variance, derivation of confidence intervals, and also a model selection procedure. The problem of prediction for new observations on a test set with missing covariate data is also tackled. Supported by a simulation study in which the method is compared to previous ones, it has proved to be computationally efficient, and has good coverage and variable selection properties. The approach is then illustrated on a dataset of severely traumatized patients from Paris hospitals by predicting the occurrence of hemorrhagic shock, a leading cause of early preventable death in severe trauma cases. The aim is to improve the current red flag procedure, a binary alert identifying patients with a high risk of severe hemorrhage. The method is implemented in the R package *misaem*.

© 2020 Elsevier B.V. All rights reserved.

1. Introduction

Missing data exist in almost all areas of empirical research. There are various reasons why it may occur, including survey non-response, unavailability of measurements, and lost data. One popular approach to handle missing values consists in modifying estimation processes so that they can be applied to incomplete data. For example, one can use the EM algorithm (Dempster et al., 1977) to obtain the maximum likelihood estimate (MLE) despite missing values, accompanied by a supplemented EM algorithm (SEM) (Meng and Rubin, 1991) or Louis' formula (Louis, 1982) for estimating the variance. This strategy is valid under missing-at-random (MAR) mechanisms (Little and Rubin, 2002; Seaman et al., 2013) in which data missingness is independent of the missing values, given the observed data. Even though this approach is perfectly suited to specific inference problems with missing values, there are few solutions or implementations available, even for simple models such as logistic regression, the focus of this paper.

One explanation is that the expectation step of the EM algorithm often involves unfeasible computations. In the framework of generalized linear models, Ibrahim et al. (1999, 2005) have suggested using a Monte Carlo EM (MCEM) algorithm (Wei and Tanner, 1990; McLachlan and Krishnan, 2008), replacing the integral by its empirical sum using Monte

* Corresponding author.

E-mail address: wei.jiang@polytechnique.edu (W. Jiang).

Carlo sampling. Ibrahim et al. (1999) also estimated the variance using a Monte Carlo version of Louis' formula, involving Gibbs sampling with an adaptive rejection sampling scheme (Gilks and Wild, 1992). However, their approach has a high computational cost and was only implemented for monotone patterns of missing values and for missing values in only two variables in a dataset.

In this paper, we develop a stochastic approximation version of the EM algorithm (SAEM) (Lavielle, 2014), based on Metropolis–Hastings sampling, to perform statistical inference for logistic regression with incomplete data, where the missing data is found anywhere in the covariates. SAEM uses a stochastic approximation procedure to estimate the conditional expectation of the complete-data likelihood, instead of generating a large number of Monte Carlo samples, which lead to an undeniable computational advantage over MCEM, which we illustrate in the simulations. In addition, SAEM allows for model selection using a criterion based on a penalized version of the observed-data likelihood. This is very useful in practice, as few methods are available to select models when there are missing values. Of those that do exist, Claeskens and Consentino (2008), Consentino and Claeskens (2011) suggested an approximation of AIC, Jiang et al. (2015) defined generalized information criteria, and – in the imputation framework – Liu et al. (2016) proposed combining penalized regression techniques with multiple imputation and stability selection. As well as aiming to maximize the MLE for observed data, Chow (1979), Yuen Fung and Wrobel (1989) also studied the linear discriminant function for logistic regression, using pairs of observed values in different columns to calculate the covariance matrix. Note that another solution is to use a Laplace approximation to compute integrals; this linearizes the likelihood function via differentiation, whereas SAEM supports likelihood-based inference without the intermediate step.

This paper proceeds as follows: In Section 2 we describe the motivation for the work: the TraumaBase¹ project involving a French multicenter prospective trauma registry. Section 3 provided the assumptions and notation used throughout the paper. In Section 4, we derive a SAEM algorithm to obtain the maximum likelihood estimate of parameters in a logistic regression model for continuous covariate data, under the MAR mechanism of missing data. Following parameter estimation, we show how to estimate the Fisher information matrix using a Monte Carlo version of Louis' formula. Section 5 describes the model selection scheme, based on a Bayesian information criterion (BIC) with missing values. In addition, we propose an approach to perform prediction for a new observation that includes missing values. Section 6 presents a simulation study where the proposed approach is compared to methods such as multiple imputation (Rubin, 2009) with respect to bias, coverage, and execution time. In Section 7, we apply the newly developed approach to predict the occurrence of hemorrhagic shock in patients with blunt trauma in the TraumaBase dataset, where it is crucial to efficiently manage missing data because the percentage of it varies from 0 to 60% depending on the variable. Predictions made using SAEM show an improvement over those made by emergency doctors. Lastly, Section 8 discusses the results and provides conclusions.

Our contribution is to give users the ability to perform logistic regression with missing values within a joint-modeling framework, one that combines computational efficiency and a sound theoretical foundation. The methods presented in this article are implemented as an R (R Core Team, 2017) package *misaem* (Jiang, 2019a), available on CRAN. The code to reproduce all simulations can be found on GitHub (Jiang, 2019b).

2. Medical emergencies

Our work is motivated by a collaboration with the TraumaBase group at APHP (Public Assistance – Hospitals of Paris), which is dedicated to the management of severely traumatized patients.

Major trauma refers to injuries that cause prolonged disability or endanger a person's life. The World Health Organization has recently reported that major trauma, such as from road accidents, interpersonal violence and falls, is a prominent source of mortality and morbidity around the world (Hay et al., 2017). In particular, major trauma is the leading cause of mortality and second cause of disability in the 16–45 age group, while hemorrhagic shock and traumatic brain injuries are the two leading causes of death.

The path of a traumatized patient involves several stages: the accident site, where care is typically provided by ambulance, transfer to an intensive-care unit for immediate intervention, followed by comprehensive care at the hospital. Using pre-hospital patient records, we aim to develop models to predict the risk of severe hemorrhage, in order to prepare an appropriate response upon arrival at a trauma center, e.g., a massive transfusion protocol and/or immediate haemostatic procedures.

Due to the highly stressful and multi-player environments involved, evidence suggests that patient management – even in mature trauma systems – often exceeds acceptable time frames (Hamada et al., 2014). In addition, discrepancies may be observed between the diagnoses made by emergency doctors in the ambulance and those made when the patient arrives at the trauma center (Hamada et al., 2015). Such discrepancies can result in poor outcomes like inadequate hemorrhage control and delayed transfusion.

To improve decision-making and patient care, 15 French trauma centers have collaborated to collect detailed high-quality clinical data from the accident site right through to the hospital. The resulting database, TraumaBase, is a multicenter prospective trauma registry that is continually updated, and now has data from more than 7000 trauma cases. The sheer quantity of collected data (more than 250 variables) makes this dataset unique in Europe. However,

¹ <http://www.traumabase.eu/>.

these data – coming from multiple sources – have high inter-center variability, not to mention the fact that a lot of data are missing, both of which problems make modeling challenging.

In this paper, we focus on performing logistic regression with missing values to help propose an innovative response to the public health challenge of major trauma.

3. Assumptions and notation

Let (y, x) be the observed data with $y = (y_i, 1 \leq i \leq n)$ an n -vector of binary responses coded as $\{0, 1\}$ and $x = (x_{ij}, 1 \leq i \leq n, 1 \leq j \leq p)$ an $n \times p$ matrix of covariates, where x_{ij} takes its values in \mathbb{R} . The logistic regression model for binary classification can be written as:

$$\mathbb{P}(y_i = 1|x_i; \beta) = \frac{\exp(\beta_0 + \sum_{j=1}^p \beta_j x_{ij})}{1 + \exp(\beta_0 + \sum_{j=1}^p \beta_j x_{ij})}, \quad i = 1, \dots, n, \quad (1)$$

where x_{i1}, \dots, x_{ip} are the covariates for individual i and $\beta_0, \beta_1, \dots, \beta_p$ unknown parameters. We adopt a probabilistic framework by assuming that $x_i = (x_{i1}, \dots, x_{ip})$ is normally distributed:

$$x_i \underset{\text{i.i.d.}}{\sim} \mathcal{N}_p(\mu, \Sigma), \quad i = 1, \dots, n.$$

Let $\theta = (\mu, \Sigma, \beta)$ be the set of parameters of the model. Then, the log-likelihood for the complete data can be written as:

$$\begin{aligned} \mathcal{LL}(\theta; x, y) &= \sum_{i=1}^n \mathcal{LL}(\theta; x_i, y_i) \\ &= \sum_{i=1}^n \left(\log(p(y_i|x_i; \beta)) + \log(p(x_i; \mu, \Sigma)) \right). \end{aligned}$$

Our main goal is to estimate the vector of parameters $\beta = (\beta_j, 0 \leq j \leq p)$ when missing values exist in the design matrix, i.e., in the matrix x . For each individual i , we note $x_{i,\text{obs}}$ the elements of x_i that are observed and $x_{i,\text{mis}}$ those that are missing. We also decompose the design matrix as $x = (x_{\text{obs}}, x_{\text{mis}})$, keeping in mind that the missing elements may differ from one individual to another.

For each individual i , we define the missing data indicator vector $M_i = (M_{ij}, 1 \leq j \leq p)$, with $M_{ij} = 1$ if x_{ij} is missing and $M_{ij} = 0$ otherwise. The matrix $M = (M_i, 1 \leq i \leq n)$ then defines the missing data pattern. The missing data mechanism is characterized by the conditional distribution of M given x and y , with parameter ϕ , i.e., $p(M_i|x_i, y_i, \phi)$. Throughout this paper, we assume a missing-at-random (MAR) mechanism, which implies that the missing values mechanism can therefore be ignored (Little and Rubin, 2002) and the maximum likelihood estimate of θ obtained by maximizing $\mathcal{LL}(\theta; y, x_{\text{obs}})$. A reminder of these concepts is given in Appendix A.1.

4. Parameter estimation by SAEM

4.1. The EM and MCEM algorithms

We aim to estimate the parameter θ of the logistic regression model by maximizing the observed log-likelihood $\mathcal{LL}(\theta; x_{\text{obs}}, y)$. Let us start with the classical EM formulation for obtaining the maximum likelihood estimator from incomplete data. Given some initial value θ_0 , iteration k updates θ_{k-1} to θ_k with the following two steps:

- **E-step:** Evaluate the quantity

$$\begin{aligned} Q_k(\theta) &= \mathbb{E}[\mathcal{LL}(\theta; x, y)|x_{\text{obs}}, y; \theta_{k-1}] \\ &= \int \mathcal{LL}(\theta; x, y) p(x_{\text{mis}}|x_{\text{obs}}, y; \theta_{k-1}) dx_{\text{mis}}. \end{aligned} \quad (2)$$

- **M-step:** Update the estimation of θ : $\theta_k = \arg \max_{\theta} Q_k(\theta)$.

Since the expectation (2) in the E-step for the logistic regression model has no explicit expression, MCEM (Wei and Tanner, 1990; Ibrahim et al., 1999) can be used. The E-step of MCEM generates several samples of missing data from the target distribution $p(x_{\text{mis}}|x_{\text{obs}}, y; \theta_{k-1})$ and replaces the expectation of the complete log-likelihood by an empirical mean. However, an accurate Monte Carlo approximation of the E-step may require a significant computational effort, as illustrated in Section 6.

4.2. The SAEM algorithm

To achieve improved computational efficiency, we can derive a SAEM algorithm (Lavielle, 2014) which replaces the E-step (2) by a stochastic approximation. Starting from an initial guess θ_0 , the k th iteration consists of three steps:

- **Simulation:** For $i = 1, 2, \dots, n$, draw $x_{i,\text{mis}}^{(k)}$ from

$$p(x_{i,\text{mis}} | x_{i,\text{obs}}, y_i; \theta_{k-1}). \quad (3)$$

- **Stochastic approximation:** Update the function Q according to

$$Q_k(\theta) = Q_{k-1}(\theta) + \gamma_k \left(\mathcal{L}(\theta; x_{\text{obs}}, x_{\text{mis}}^{(k)}, y) - Q_{k-1}(\theta) \right), \quad (4)$$

where (γ_k) is a non-increasing sequence of positive numbers.

- **Maximization:** Update the estimation of θ :

$$\theta_k = \arg \max_{\theta} Q_k(\theta).$$

The choice of the sequence (γ_k) in (4) is important for ensuring the almost sure convergence of SAEM to a maximum of the observed likelihood (Delyon et al., 1999). We will see in Section 6 that, in our case, very good convergence is obtained using $\gamma_k = 1$ during the first iterations, followed by a sequence that decreases as $1/k$.

4.3. Metropolis–Hastings sampling

In the logistic regression case, the unobserved data cannot in general be drawn exactly from the conditional distribution (3), which depends on an integral that is not calculable in closed form. One solution is to use a Metropolis–Hastings (MH) algorithm, which consists of constructing a Markov chain that has the target distribution as its stationary distribution. The states of the chain after S iterations are then used as a sample from the target distribution. To define a proposal distribution for the MH algorithm, we observe that the target distribution (3) can be factorized as follows:

$$p(x_{i,\text{mis}} | x_{i,\text{obs}}, y_i; \theta) \propto p(y_i | x_i; \beta) p(x_{i,\text{mis}} | x_{i,\text{obs}}; \mu, \Sigma).$$

We select the proposal distribution as the second term $p(x_{i,\text{mis}} | x_{i,\text{obs}}, \mu, \Sigma)$, which is normally distributed:

$$x_{i,\text{mis}} | x_{i,\text{obs}} \sim \mathcal{N}_p(\mu_i, \Sigma_i), \quad (5)$$

where

$$\begin{aligned} \mu_i &= \mu_{i,\text{mis}} + \Sigma_{i,\text{mis},\text{obs}} \Sigma_{i,\text{obs},\text{obs}}^{-1} (x_{i,\text{obs}} - \mu_{i,\text{obs}}), \\ \Sigma_i &= \Sigma_{i,\text{mis},\text{mis}} - \Sigma_{i,\text{mis},\text{obs}} \Sigma_{i,\text{obs},\text{obs}}^{-1} \Sigma_{i,\text{obs},\text{mis}}, \end{aligned}$$

with $\mu_{i,\text{mis}}$ (resp. $\mu_{i,\text{obs}}$) the missing (resp. observed) elements of μ for individual i . The covariance matrix Σ is decomposed in the same way. The MH algorithm is described further in Appendix A.2.

4.4. Observed Fisher information

After computing the MLE $\hat{\theta}_{\text{ML}}$ with SAEM, we estimate its variance. To do so, we can use the observed Fisher information matrix (FIM): $\mathcal{I}(\theta) = -\frac{\partial^2 \mathcal{L}(\theta; x_{\text{obs}}, y)}{\partial \theta \partial \theta^T}$. According to Louis' formula (Louis, 1982), we have:

$$\begin{aligned} \mathcal{I}(\theta) &= -\mathbb{E} \left(\frac{\partial^2 \mathcal{L}(\theta; x, y)}{\partial \theta \partial \theta^T} | x_{\text{obs}}, y; \theta \right) \\ &\quad - \mathbb{E} \left(\frac{\partial \mathcal{L}(\theta; x, y)}{\partial \theta} \frac{\partial \mathcal{L}(\theta; x, y)^T}{\partial \theta} | x_{\text{obs}}, y; \theta \right) \\ &\quad + \mathbb{E} \left(\frac{\partial \mathcal{L}(\theta; x, y)}{\partial \theta} | x_{\text{obs}}, y; \theta \right) \mathbb{E} \left(\frac{\partial \mathcal{L}(\theta; x, y)}{\partial \theta} | x_{\text{obs}}, y; \theta \right)^T. \end{aligned}$$

The observed FIM can therefore be expressed in terms of conditional expectations, which can also be approximated using a Monte Carlo procedure. More precisely, given S samples $(x_{i,\text{mis}}^{(s)}, 1 \leq i \leq n, 1 \leq s \leq S)$ of the missing data drawn from

the conditional distribution (3), the observed FIM can be estimated as $\hat{\mathcal{I}}_S(\hat{\theta}) = \sum_{i=1}^n -(D_i + G_i - \Delta_i \Delta_i^T)$, where

$$\begin{aligned}\Delta_i &= \frac{1}{S} \sum_{s=1}^S \frac{\partial \mathcal{L}(\hat{\theta}; x_{i,\text{mis}}^{(s)}, x_{i,\text{obs}}, y_i)}{\partial \theta}, \\ D_i &= \frac{1}{S} \sum_{s=1}^S \frac{\partial^2 \mathcal{L}(\hat{\theta}; x_{i,\text{mis}}^{(s)}, x_{i,\text{obs}}, y_i)}{\partial \theta \partial \theta^T}, \\ G_i &= \frac{1}{S} \sum_{s=1}^S \left(\frac{\partial \mathcal{L}(\hat{\theta}; x_{i,\text{mis}}^{(s)}, x_{i,\text{obs}}, y_i)}{\partial \theta} \right) \left(\frac{\partial \mathcal{L}(\hat{\theta}; x_{i,\text{mis}}^{(s)}, x_{i,\text{obs}}, y_i)}{\partial \theta} \right)^T.\end{aligned}$$

Here, the gradient and the Hessian matrix can be computed in closed form. The procedure for calculating the observed information matrix is described in [Appendix A.3](#).

5. Model selection and prediction

5.1. Information criteria

In order to compare different possible covariate models, we can consider penalized likelihood criteria such as BIC. For a given model \mathcal{M} and an estimated parameter $\hat{\theta}_{\mathcal{M}}$, the BIC is defined as:

$$\text{BIC}(\mathcal{M}) = -2\mathcal{L}(\hat{\theta}_{\mathcal{M}}; x_{\text{obs}}, y) + \log(n)d(\mathcal{M}),$$

where $d(\mathcal{M})$ is the number of estimated parameters in a model \mathcal{M} . The distribution of the complete set of covariates $(x_{ij}, 1 \leq i \leq n, 1 \leq j \leq p)$ does not depend on the regression model used for modeling the binary outcomes $(y_i, 1 \leq i \leq n)$; we assume the same normal distribution $\mathcal{N}_p(\mu, \Sigma)$ for all regression models. Thus, the difference between the numbers $d(\mathcal{M})$ of estimated parameters in two models is equivalent to the difference between the numbers of their non-zero coefficients in β . Note that, unlike the approach we suggest, existing methods ([Claeskens and Consentino, 2008](#); [Consentino and Claeskens, 2011](#)) use an approximation of the Akaike information criterion (AIC) without estimating the observed likelihood.

5.2. Observed log-likelihood

For a given model and parameter θ , the observed log-likelihood is, by definition:

$$\mathcal{L}(\theta; x_{\text{obs}}, y) = \sum_{i=1}^n \log(p(y_i, x_{i,\text{obs}}; \theta)).$$

With missing data, the density $p(y_i, x_{i,\text{obs}}; \theta)$ cannot in general be computed in closed form. We suggest approximating it using an importance [sampling Monte Carlo approach](#). Let g_i be the density function of the normal distribution defined in (5). Then,

$$\begin{aligned}p(y_i, x_{i,\text{obs}}; \theta) &= \int p(y_i, x_{i,\text{obs}} | x_{i,\text{mis}}; \theta) p(x_{i,\text{mis}}; \theta) dx_{i,\text{mis}} \\ &= \int p(y_i, x_{i,\text{obs}} | x_{i,\text{mis}}; \theta) \frac{p(x_{i,\text{mis}}; \theta)}{g_i(x_{i,\text{mis}})} g_i(x_{i,\text{mis}}) dx_{i,\text{mis}} \\ &= \mathbb{E}_{g_i} \left(p(y_i, x_{i,\text{obs}} | x_{i,\text{mis}}; \theta) \frac{p(x_{i,\text{mis}}; \theta)}{g_i(x_{i,\text{mis}})} \right).\end{aligned}$$

Consequently, if we draw M samples from the proposal distribution (5):

$$x_{i,\text{mis}}^{(s)} \underset{\text{i.i.d.}}{\sim} \mathcal{N}(\mu_i, \Sigma_i), \quad m = 1, 2, \dots, S,$$

we can estimate $p(y_i, x_{i,\text{obs}}; \theta)$ by:

$$\hat{p}(y_i, x_{i,\text{obs}}; \theta) = \frac{1}{S} \sum_{m=1}^S p(y_i, x_{i,\text{obs}} | x_{i,\text{mis}}^{(s)}; \theta) \frac{p(x_{i,\text{mis}}^{(s)}; \theta)}{g_i(x_{i,\text{mis}}^{(s)})},$$

and derive an estimate of the observed log-likelihood $\mathcal{L}(\theta; x_{\text{obs}}, y)$.

5.3. Prediction on a test set with missing values

In supervised learning, after fitting a model using a training set, a natural step is to evaluate the prediction performance, which can be done with a test set. Assuming $\tilde{x} = (\tilde{x}_{\text{obs}}, \tilde{x}_{\text{mis}})$ is an observation in the test set, we want to predict the binary response \tilde{y} . One important point is that test set also contains missing values, since the training set and the test set have the same distribution (i.e., the distribution of covariates and the distribution of missingness). Therefore, we cannot directly apply the fitted model (which uses p coefficients) to predict \tilde{y} from an incomplete observation of the test set \tilde{x} .

Our framework offers a natural way to tackle this issue by marginalizing out missing covariates given the observed data. More precisely, with S Monte Carlo samples

$$(\tilde{x}_{\text{mis}}^{(s)}, 1 \leq s \leq S) \sim p(\tilde{x}_{\text{mis}} | \tilde{x}_{\text{obs}}),$$

we estimate directly the response by maximizing its distribution marginalized over missing data given the observed ones:

$$\begin{aligned} \hat{y} &= \arg \max_{\tilde{y}} p(y | \tilde{x}_{\text{obs}}) = \arg \max_{\tilde{y}} \int p(y | \tilde{x}) p(\tilde{x}_{\text{mis}} | \tilde{x}_{\text{obs}}) d\tilde{x}_{\text{mis}} \\ &= \arg \max_{\tilde{y}} \mathbb{E}_{p_{\tilde{x}_{\text{mis}} | \tilde{x}_{\text{obs}}}} p(\tilde{y} | \tilde{x}) \\ &= \arg \max_{\tilde{y}} \sum_{s=1}^S p(\tilde{y} | \tilde{x}_{\text{obs}}, \tilde{x}_{\text{mis}}^{(s)}). \end{aligned}$$

Note that in the literature there are very few solutions for dealing with missing values in a test set. In Section 7.2, we compare the proposed approach with other methods used in practice, which are based on imputation of the test set.

6. Simulation study

6.1. Simulation settings

We first generated a design matrix x of size $n = 1000 \times p = 5$ by drawing each observation from a multivariate normal distribution $\mathcal{N}(\mu, \Sigma)$. Then, we generated the response according to the logistic regression model (1). We considered as the true parameter values: $\beta = (-0.2, 0.5, -0.3, 1, 0, -0.6)$, $\mu = (1, 2, 3, 4, 5)$, and $\Sigma = \text{diag}(\sigma)C\text{diag}(\sigma)$, where σ is the vector of standard deviations $\sigma = (1, 2, 3, 4, 5)$, and C the correlation matrix

$$C = \begin{bmatrix} 1 & 0.8 & 0 & 0 & 0 \\ 0.8 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0.3 & 0.6 \\ 0 & 0 & 0.3 & 1 & 0.7 \\ 0 & 0 & 0.6 & 0.7 & 1 \end{bmatrix}. \quad (6)$$

Before generating missing values, we performed classical logistic regression on the complete dataset, the results (ROC curve) are provided in Appendix A.4. We then randomly introduced 10% missing values in the covariates, initially with a missing-completely-at-random (MCAR) mechanism, where each entry has the same probability of being observed.

6.2. The behavior of SAEM

The algorithm was initialized with the parameters obtained after mean imputation, i.e., where missing values of a given variable were replaced by the unconditional mean calculated on the available cases, and then logistic regression was applied to the completed data. For the non-increasing sequence (γ_k) in the stochastic approximation step of SAEM, we chose $\gamma_k = 1$ during the first k_1 iterations in order to converge quickly to a neighborhood of the MLE, and from k_1 iterations on, we set $\gamma_k = (k - k_1)^{-\tau}$ to ensure the almost sure convergence of SAEM. In order to study the effect of the sequence of stepsizes (γ_k) , we fixed the value of $k_1 = 50$ and used $\tau = (0.6, 0.8, 1)$ during the next 450 iterations. Representative plots of the convergence of SAEM for the coefficient β_1 , obtained from four simulated datasets, are shown in Fig. 1. For each given simulation, the three sequences of estimates converged to the same solution, but for larger τ , SAEM converged faster and fluctuated less. We therefore use $\tau = 1$ in the following.

6.3. Comparison with other methods

We ran 1000 simulations and compared SAEM to several other existing methods, initially in terms of estimation errors for the parameters. We mainly focused on (i) the complete case (CC) method, i.e., all rows containing at least one unobserved data value were removed, and (ii) multiple imputation by chained equations (*mice*) with Rubin's combining rules (van Buuren and Groothuis-Oudshoorn, 2011). More precisely, missing values are imputed successively by a series of regression models, where each variable with missing data is modeled conditional upon the other variables. For instance, linear regression is used to model continuous variables and binary variables are modeled using logistic regression. More




Fig. 1. Convergence plots for β_1 obtained with three different values of τ (0.6, 0.8, 1.0). Each color represents one simulation. The true value of β_1 is 0.5.

Table 1

Coverage (%) for $n = 10\,000$, correlation C and 10% MCAR, calculated over 1000 simulations. Bold indicates under-coverage. Inside the parentheses is the average length of corresponding confidence interval over 1000 simulations (multiplied by 100).

Parameter	no NA	CC	mice	SAEM
β_0	95.2 (21.36)	94.4 (27.82)	95.2 (22.70)	94.9 (22.48)
β_1	96.0 (18.92)	94.7 (24.65)	93.9 (21.77)	95.1 (21.51)
β_2	95.5 (9.53)	94.6 (12.41)	94.0 (10.97)	94.3 (10.83)
β_3	94.9 (8.17)	94.3 (10.66)	86.5 (9.03)	94.7 (9.03)
β_4	94.6 (4.00)	94.2 (5.21)	96.2 (4.49)	95.4 (4.42)
β_5	95.9 (5.52)	94.4 (7.19)	89.6 (6.20)	94.7 (6.17)

details can be found in [van Buuren and Groothuis-Oudshoorn \(2011\)](#). Finally, we used the dataset without missing values ("no NA") as a reference, with parameters estimated using the Newton–Raphson algorithm. We varied the number of observations: $n = 200, 1000, 10\,000$, the missing value mechanism: MCAR or MAR, the percentage of missing values: 10% or 30%, and the correlation structure, either using C given by (6), or an orthogonal design.

[Fig. 2](#) (top) displays the distribution of the estimates of β_3 for $n = 1000$ and $n = 10\,000$ under MCAR, with the correlation between covariates given by (6). Simulation results for $n = 200$ are presented in the supplementary materials ([Jiang, 2019c](#)). This plot is representative of the results obtained with the other components of β . As expected, larger samples yielded less variability. Moreover, we observe that in both cases, the estimation obtained by *mice* could be biased, whereas SAEM provided unbiased estimates with small variances. [Fig. 2](#) (bottom) shows the empirical distribution of the estimated standard error of $\hat{\beta}_3$. For SAEM it was calculated using the observed Fisher information as described in Section 4.4. With larger n , not only the estimated standard errors – but also variance in the estimation – clearly decreased for all methods. In the case where $n = 1000$, SAEM and *mice* slightly overestimated the standard error, while CC underestimated it, on average. Globally, SAEM provided the best results; compared with *mice*, it gave a similar estimate of the standard error, on average, but with much less variance.

[Table 1](#) shows the coverage probability of the confidence interval for all parameters and inside the parentheses is the average length of the corresponding confidence interval. We would expect coverage of 95%, corresponding to the nominal 95% level. The simulation margin of error for the coverage results is 1.35%. SAEM had between 94.3% and 95.4% coverage, while *mice* struggled for certain parameters: the coverage rates for two estimates were 89.6% and 86.5%, significantly below the nominal level. Even though CC showed reasonable results in terms of coverage, the widths of its confidence intervals were still too large. Simulations with smaller sample sizes gave similar results—see supplementary materials ([Jiang, 2019c](#)) for $n = 200$.

Lastly, [Table 2](#) highlights large differences between the methods in terms of execution time. As an aside, we also implemented the MCEM algorithm ([Ibrahim et al., 1999](#)) using adaptive rejection sampling; even with a very small sample




Fig. 2. Top: Empirical distribution of the bias of $\hat{\beta}_3$. Bottom: Distribution of the estimated standard errors of $\hat{\beta}_3$. For each method, the red point corresponds to the empirical standard deviation of $\hat{\beta}_3$ calculated over the 1000 simulations. Results shown are for 10% MCAR and correlation C.

Table 2

Comparison of execution times between no NA, MCEM, *mice*, and SAEM with correlation C and 10% MCAR, for $n = 200$ and $n = 1000$, calculated over 1000 simulations.

Execution time (seconds) for one simulation	no NA	MCEM	<i>mice</i>	SAEM
$n = 1000$				
min	2.87×10^{-3}	492	0.64	9.96
mean	4.65×10^{-3}	773	0.70	13.50
max	43.50×10^{-3}	1077	0.76	16.79
$n = 200$				
min	1.26×10^{-3}	67.91	0.24	2.64
mean	2.32×10^{-3}	291.47	0.28	3.91
max	21.53×10^{-3}	1003	0.48	6.04

size of $n = 200$, MCEM took 5 min per simulation on average. In contrast, multiple imputation took less than 1 s per simulation, and SAEM less than 10 s, which remains reasonable. However, the bias and standard errors for the SAEM and MCEM estimates were quite similar—see supplementary materials (Jiang, 2019c). Due to the prohibitive execution time required, for larger sample sizes we did not compare MCEM with the other methods.

Results obtained for independent covariates are presented in Fig. 3 (right), for estimation in the orthogonal design case. SAEM was slightly biased since it estimated non-zero terms for the covariance, but still outperformed CC and *mice*.

6.4. Extended simulations

Missing-at-random mechanisms. We first simulated the pattern of missingness as a binary vector $\eta = (\eta_1, \eta_2, \dots, \eta_p)$ from the Bernoulli distribution, where $\eta_j = 0$ indicates that the corresponding variable x_j can be missing while $\eta_j = 1$ indicates it is always observed. Then the probability of having missing data in one variable is calculated by a logistic regression model. For example in our case with the realizations of $\eta = (1, 0, 1, 0, 0)$, the probability that covariates (x_2, x_4, x_5) are




Fig. 3. Empirical distribution of the estimates of β_3 obtained under MCAR, with $n = 10\,000$ and 10% missing values. Left: the covariates are correlated; right: no correlation between covariates.




Fig. 4. Empirical distribution of the bias of $\hat{\beta}_3$ obtained for misspecified models under MCAR, with $n = 1000$. Left: Student's distribution with $\nu = 5$ degrees of freedom; right: Gaussian mixture model.

missing is calculated by a logistic regression model conditional on x_1 and x_3 . The weights in the linear combination of x_1 and x_3 have an effect on the proportion of missingness. We introduced 10% missing values into the covariates using an MAR mechanism. The results presented in [Appendix A.5](#) are – as expected – similar to those obtained under MCAR, and show that the parameters are estimated without bias.

Robustness to the normal assumption for covariates. First we generated a design matrix of size $n = 1000 \times p = 5$ by drawing each observation from a multivariate Student distribution $t_v(\mu, \Sigma)$ with $\nu = 5$ or $\nu = 20$ degrees of freedom, and (μ, Σ) the same as those in the normal distribution in [Section 6.1](#). Then, we considered the Gaussian mixture model case by generating half of the samples from $\mathcal{N}(\mu_1, \Sigma)$ and the other half from $\mathcal{N}(\mu_2, \Sigma)$, where $\mu_1 = (1, 2, 3, 4, 5)$ and $\mu_2 = (1, 1, 1, 1, 1)$, with the same Σ as previously. Then, we generated the response according to the same logistic regression model as in [Section 6.1](#), and considered either MCAR or MAR mechanisms.

[Fig. 4](#) illustrates the estimation bias of the parameter β_3 , and [Appendix A.6](#) shows the coverage for all parameters, with the average length of the corresponding confidence interval in parentheses. This experiment shows that the estimation bias for regression coefficients with the proposed method – even based on normal assumptions – is robust to such model misspecification. Indeed, the bias may increase when covariates do not exactly follow a normal distribution, but the increase is negligible compared to the bias of imputation-based methods. We also observe only a small level of under-coverage as compared to *mice*, and a more reasonable length of confidence interval as compared to CC.

Varying the percentage of missing values. When the percentage of missing values increases, variability in the results increases, but the suggested method still provides satisfactory results—see supplementary materials ([Jiang, 2019c](#)).

Varying the separability of the classes. When the classes are well-separated, SAEM can exhibit bias and large variance, as illustrated in the supplementary materials ([Jiang, 2019c](#)). However, logistic regression without missing values also encounters difficulties.

Table 3

For data with or without correlations, the percentage of times that each criterion selects the correct true model (C), overfits (O), or underfits (U).

Criterion	Non-Correlated			Correlated		
	C	O	U	C	O	U
BIC_{obs}	92	3	5	94	2	4
BIC_{orig}	96	2	2	93	0	7
BIC_{cc}	79	1	20	91	0	9




Fig. 5. Comparisons of the empirical distribution of the AUC, Brier score, and logarithmic score obtained on the test set for the proposed SAEM without imputation method, impMean, impPCA, and mice, over 100 simulations.

In summary, not only did these simulations show that SAEM leads to estimators with limited bias, but also that we obtained accurate inference by taking into account the additional variance due to missing data.

6.5. Model selection

To look at the performance of the method in terms of model selection, we considered the same simulation scenarios as in Section 6.1, with some parameters set to zero. We now describe the results for the case where all parameters in β are zero except $\beta_0 = -0.2$, $\beta_1 = 0.5$, $\beta_3 = 1$, and $\beta_5 = -0.6$. We compared the BIC_{obs} based on the observed log-likelihood, as described in Section 5, to those based on the complete cases BIC_{cc} and obtained from the original complete data BIC_{orig} .

Table 3 shows, with or without correlation between covariates, the percentage of cases where each criterion selects the true model (C), overfits (O)—i.e., selects more variables than there were, or underfits (U)—i.e., selects less variables than there were. In the case where the variables were correlated, the correlation matrix was the same as in Section 6.1. These results are representative of those obtained in the other simulation settings.

6.6. Predictions for a test set with missing values

To evaluate the prediction performance on a test set with missing values, we considered the same simulation scenarios for the training set as in Section 6.1 with sample size 1000×5 . We also generated a test set of size 100×5 . We compared the approach described in Section 5.3 with imputation methods. More precisely, we considered single imputation methods on the training set, followed by classical logistic regression and variable selection by BIC on the imputed dataset. The single imputation methods included (i) imputation by the mean (impMean) (ii) imputation by PCA (impPCA) (Josse and Husson, 2016), which is based on a low-rank assumption of the data matrix to impute. In addition, we considered multiple imputation using mice. Note that Hentges and Dunsmore (1998) showed in a simulation study that imputation methods for MCAR data can perform well when the aim of logistic regression is prediction. For all of the imputation methods, we also imputed the test set independently and then applied the model that had been selected on the training set. Note that this would be a limitation if there was only one individual in the test set, whereas our method does not encounter this issue.

We compared all of these approaches in terms of classical criteria to evaluate the predicted probabilities from the logistic regression. Criteria included AUC (area under the ROC curve), the Brier score (Brier, 1950) and the logarithmic score (Good, 1952). Fig. 5 shows that on average, marginalizing over the distribution of missing values gave the best performance: the largest AUCs and logarithmic scores, and the smallest Brier scores.




Fig. 6. The factor map of the variables from PCA.

7. Risk of severe hemorrhage in the TraumaBase context

The fundamental goal of our work is to accelerate and simplify the detection of patients presenting with hemorrhagic shock due to blunt trauma in order to speed up management of this, the most preventable cause of death in major trauma cases. Optimized organization is essential to control blood loss as quickly as possible and reduce mortality.

7.1. Details on the dataset

This study has used data collected from the TraumaBase[®] trauma registry, which collates data from six trauma centers in the Ile-de-France region (Paris area) in France. The centers progressively joined TraumaBase between January 2011 and June 2015. Since then, data collection has been systematic at these centers. The database is structured in such a way that the algorithm is implemented within it to provide consistency and coherence, with data monitoring performed by a central administrator. Sociodemographic, clinical, biological and therapeutic data (from the pre-hospital phase to discharge, if hospitalized) are systematically recorded for all trauma patients, and all patients transported to the emergency rooms of participating centers are included in the registry. As a result, there were 7495 individuals logged in the trauma data that we investigated, included from January 2011 to March 2016, with ages ranging from 12 to 96. The TraumaBase group decided to focus on patients with blunt trauma so as to be able to compare results with existing prediction rules. Patients with pre-hospital cardiac arrest, penetrating trauma, and missing pre-hospital data, were excluded. This led to 5162 patients being retained in the dataset. Based on clinical experience, 16 influential quantitative measurements were included. Detailed descriptions of these and their histograms are shown in [Appendix A.7](#). These variables were chosen because they were all available to the pre-hospital team, and therefore could be used in real situations.

There was strong collinearity between variables, as can be seen in the variables' PCA factor map (obtained by running an EM-PCA algorithm ([Josse and Husson, 2016](#)) which performs PCA with missing values) in [Fig. 6](#), in particular between the minimum systolic (PAS.min) and diastolic blood pressure (PAD.min). Based on expert advice, the recoded variables, SD.min and SD.SMUR ($SD.min = PAS.min - PAD.min$; $SD.SMUR = PAS.SMUR - PAD.SMUR$) were used since they have more clinical significance ([Hamada et al., 2018](#)). Thus, we had 14 variables to predict hemorrhagic shock.

[Fig. 7](#) shows the percentage of missingness per variable, varying from 0 to 60%, which demonstrates the importance of taking appropriate account of missing data. Even though there may be many reasons why missingness occurred, overall considering them all to be MAR remains a plausible assumption. For instance, FC.SMUR (heart rate) and SD.SMUR (the pulse pressure measured when the ambulance arrives at the accident site) contain many missing values because doctors collected these data during transportation. However, many other medical institutes and scientific publications use measurements on arrival at the accident scene. Consequently, doctors decided to record these as well, but this occurred after TraumaBase was set up.

We first applied SAEM for logistic regression with all 14 predictors and for the whole dataset. The estimation obtained by SAEM was broadly similar to that obtained by multiple imputation. Next, we used the model selection procedure




Fig. 7. Percentage of missing values in each variable.




Fig. 8. The observations' PCA factor map. Red points are hemorrhagic shock patients, and black points those who did not have hemorrhagic shock. Patient number 3302 (circled in blue) has an incorrectly-calculated BMI.

described in Section 5. There were two observations that led to a very small value for the log-likelihood. Upon closer inspection, we found that for patient number 3302, the BMI was obtained using an incorrect calculation, and for patient number 1144, the weight (200 kg) and height (100 cm) values were likely to be incorrect. Hence, the observed log-likelihood also helped us to identify undetected outliers. In the observations' PCA factor map shown in Fig. 8, patient number 3302 (circled in blue) is one of the outliers.

7.2. Predictive performance

We divided the dataset into training and test sets. The training set contained a random selection of 70% of the observations, and the test set contained the remaining 30%. In the training set, we selected a model with the approach

Table 4Estimation of β and its standard errors obtained by SAEM, using BIC for model selection.

Variables	Estimate (standard errors)
(Intercept)	−0.12 (0.64)
Age	0.017 (0.0037)
Glasgow.moteur	−0.22 (0.040)
FC.max	0.024 (0.0028)
Hemocue.init	−0.26 (0.033)
RT.cristalloides	0.00088 (0.00011)
RT.colloides	0.0018 (0.00023)
SD.min	−0.027 (0.0055)
SD.SMUR	−0.018 (0.0061)

Table 5

Confusion matrix for predictions on the test set.

		Predicted outcome	
		1	0
Observed value	1	True Positive (102)	False Negative (46)
	0	False Positive (146)	True Negative (1254)

suggested in Section 5, and used forward selection, resulting in a model with 8 variables. The estimates of parameters and their standard errors are shown in Table 4.

The TraumaBase medical team indicated to us that the signs of the coefficients were in agreement with their prior intuition: all things being equal, (a) Older people are more likely to have hemorrhagic shock; (b) A low Glasgow score implies little or no motor response, which often is the case for hemorrhagic shock patients; (c) A typical sign of hemorrhagic shock is rapid heart rate; (d) The more a patient bleeds, the lower their hemoglobin is and more blood must be transfused. It is then more likely they will end up with hemorrhagic shock; (e) Therapy involving two types of volume expanders, cristalloides and colloides, can be conducted to treat hemorrhagic shock; (f) If an extremely low pulse pressure is observed, the cause may be a low stroke volume, which is usually the case in hemorrhagic shock.

Next, we assessed the prediction quality on the test set with the usual metrics based on the confusion matrix (false positive rate, false negative rate, etc.). We need to ensure that the cost of a false negative is much more than that of a false positive, as non-recognition of a potential hemorrhagic shock leads to a higher risk of patient mortality. With this in mind, we define the validation error on the test set as:

$$l(\hat{y}, y) = \frac{1}{n} \sum_{i=1}^n w_0 \mathbb{1}_{\{y_i=1, \hat{y}_i=0\}} + w_1 \mathbb{1}_{\{y_i=0, \hat{y}_i=1\}} \quad (7)$$

where w_0 and w_1 are user defined weight for the cost of a false negative and false positive respectively, with $w_0 + w_1 = 1$. In this way, we can choose a threshold for the logistic regression by giving values for w_0 and w_1 . For instance, we chose $\frac{w_0}{w_1} = 5$, i.e., a false negative is five times more costly than a false positive. This cost function was chosen after discussions with experts. Note that the test set was also incomplete, so we used the strategy described in Section 5.3 to perform prediction. The confusion matrix of the predictive performance on the test set is shown in Table 5. The associated ROC curve is shown in Fig. 9, which has an AUC of 0.88.

7.3. Comparison with other approaches

Next, we compared the proposed method to other approaches. Similar to in Section 7.2, we considered single imputation methods followed by classical logistic regression and variable selection on the imputed training dataset, such as single imputation by PCA (impPCA) (Josse and Husson, 2016), imputation by Random Forest (missForest) (Stekhoven and Buehlmann, 2012), and mean imputation (impMean). We also compared the logistic regression model with other prediction models such as Random Forest (predRF) and SVM (predSVM), both applied on the Random Forest-imputed (Stekhoven and Buehlmann, 2012) dataset. We also considered mice: we applied logistic regression with a classical forward selection method, with the BIC calculated on each imputed dataset. However, note that there is no straightforward solution for




Fig. 9. ROC curve of the test set predictions.




Fig. 10. Empirical distribution of the prediction errors of different methods over 15 random splits of the TraumaBase data.

combining multiple imputation and variable selection; we followed the empirical approach suggested in Wood et al. (2008) where they select variables that appear in at least half of the models selected in each imputed dataset.

We also considered three rules used by doctors to predict hemorrhagic shock: (i) Doctors' prediction (doctor): the prediction recorded in TraumaBase. This showed whether the doctor considered the patient to be at risk of hemorrhagic shock; (ii) The assessment of blood consumption score (ABC): this is an examination usually performed when the patient arrives at the trauma center. As such, the score is not exactly pre-hospital but can be computed very early in a hospitalization; (iii) the trauma associated severe hemorrhage score (TASH): this score was also designed for hemorrhage detection, but at a later time-point since it uses some values that are only available after laboratory tests and radiography.

Fig. 10 compares the methods in terms of their validation error (7). The splitting of data (into training and test sets) was repeated 15 times and we fixed the threshold such that the cost of a false negative was five times that of a false positive, i.e., $\frac{w_0}{w_1} = 5$. On average, SAEM performed well and with low variability between trials, while all of the imputation methods performed similarly to each other even naive mean imputation. In addition, other prediction methods (Random Forest and SVM) did not give smaller errors on the test sets than the logistic regression models. Lastly, the rules used by doctors, even those using more information than pre-hospital data, did not perform as well as SAEM. Appendix A.8 gives further details on classical criteria (AUC, sensitivity, specificity, accuracy and precision) to compare the predictive




Fig. 11. Average prediction errors of different methods as a function of the cost ratio $\{\frac{w_0}{w_1} \mid \frac{w_0}{w_1} > 1\}$ taken over 15 random splits of the TraumaBase data.

performance of the methods. The SAEM approach performed well on average, and particularly well for sensitivity, i.e., it rarely misdiagnosed hemorrhagic shock patients, which gels well with the clinical needs of emergency doctors.

More generally, without defining a specific threshold, we show in Fig. 11 the average predictive loss over 15 replicates as a function of the cost ratio $\{\frac{w_0}{w_1} \mid \frac{w_0}{w_1} > 1\}$ for all methods. SAEM had a small error on the test sets given the value of $\frac{w_0}{w_1}$, especially when we increased the cost of false negatives. Note that the errors for the doctors' rules and ABC increased as a function of the cost importance $\frac{w_0}{w_1}$, which means that these rules are more conservative than SAEM is, which may be problematic in this setting. Also, missForest had excellent predictive capabilities which is consistent with the results of Josse et al. (2019). However, it is difficult to interpret the results from random forest in terms of selected variables, which is often crucial for emergency doctors.

Note that even if our proposed methodology is based on the assumption of normally distributed covariates, its performance is better than the predictions made by widely-used medical criterion in terms of prediction error. Further discussion on the normal assumption is provided in Appendix A.7. In addition, it should be noted that the proposed methodology can be extended to other assumptions about the joint distribution of covariates, such as a mixture of distributions.

In summary, based on the TraumaBase application and comparisons with other methods, we have demonstrated that this new approach has the ability to outperform existing popular methods that deal with missing data.

8. Discussion

In this paper, we have developed a comprehensive joint-modeling framework for logistic regression with missing values. The method is implemented in the R package *misaem*. The experiments we have performed indicate that this method is computationally efficient and easy to implement. In addition, compared with multiple imputation – especially in the case of correlation between variables – estimation using SAEM is less biased than other methods and generally leads to interval-estimate coverage that is close to the nominal level. Based on our algorithm, model selection with BIC and missing data can be performed in a natural way. In view of the results reported in this article, we have been invited by emergency-room doctors in one of the TraumaBase centers to implement the missing-data methodology outlined here in a prospective study to evaluate its performance in a real-time clinical setting.

Paths for possible future research include further developing the method to handle both quantitative and categorical data. This paper focused on making inference with missing values, but we have also suggested a method to predict from a test set with missing values. More work could be done in the direction of supervised learning with missing values, especially when we want to better estimate the variance of predictions. Extensions of the methods of Schafer and Schenker (2000) could be considered. In addition, in the TraumaBase dataset, it would be reasonable to expect to have both MAR and missing-not-at-random (MNAR) values. MNAR means that missingness is related to the missing values themselves,

and therefore a more correct methodology would require incorporating models for missing data mechanisms. As a final note, our proposed method may be quite useful in a causal inference framework, especially for propensity score analysis, which estimates the effect of a treatment, policy, or other intervention. Indeed, inverse probability weighting methods (IPW) are often performed with logistic regression, and the proposed method offers a potential solution for times where there are missing values in the covariates.

Acknowledgments

Wei Jiang was supported by grants from Region Ile-de-France: <https://www.dim-mathinnov.fr>. The authors are thankful for fruitful discussion with Kevin Bleakley and Antoine Ogier.

Appendix A

A.1. Missingness mechanisms

Missing-completely-at-random (MCAR) means that there is no relationship between the missingness of the data and any values, observed or missing. In other words, MCAR means:

$$p(M_i|y, x_i, \phi) = p(M_i|\phi).$$

Missing-at-random (MAR) means that the probability to have missing values may depend on the observed data, but not on the missing data. We must carefully define what this means in our case by decomposing the data x_i into a subset $x_i^{(\text{mis})}$ of data that “can be missing”, and a subset $x_i^{(\text{obs})}$ of data that “cannot be missing”, i.e., that is always observed. Then, the observed data $x_{i,\text{obs}}$ necessarily includes the data that can be observed $x_i^{(\text{obs})}$, while the data that can be missing $x_i^{(\text{mis})}$ includes the missing data $x_{i,\text{mis}}$. Thus, the MAR assumption implies that, for each individual i ,

$$\begin{aligned} p(M_i|y_i, x_i; \phi) &= p(M_i|y_i, x_i^{(\text{obs})}; \phi) \\ &= p(M_i|y_i, x_{i,\text{obs}}; \phi). \end{aligned}$$

The MAR assumption implies that the observed likelihood can be maximized and the distribution of M can be ignored (Little and Rubin, 2002). Indeed,

$$\begin{aligned} \mathcal{L}(\theta, \phi; y, x_{\text{obs}}, M) &= p(y, x_{\text{obs}}, M; \theta, \phi) \\ &= \prod_{i=1}^n p(y_i, x_{i,\text{obs}}, M_i; \theta, \phi) \\ &= \prod_{i=1}^n \int p(y_i, x_i, M_i; \theta, \phi) dx_{i,\text{mis}} \\ &= \prod_{i=1}^n \int p(y_i, x_i; \theta) p(M_i|y_i, x_i; \phi) dx_{i,\text{mis}} \\ &= \prod_{i=1}^n \int p(y_i, x_i; \theta) p(M_i|y_i, x_{i,\text{obs}}; \phi) dx_{i,\text{mis}} \\ &= \prod_{i=1}^n p(M_i|y_i, x_{i,\text{obs}}; \phi) \times \prod_{i=1}^n \int p(y_i, x_i; \theta) dx_{i,\text{mis}} \\ &= p(M|y, x_{\text{obs}}; \phi) \times p(y, x_{\text{obs}}; \theta) \\ &= p(M|y, x^{(\text{obs})}; \phi) \times p(y, x_{\text{obs}}; \theta). \end{aligned}$$

Therefore, to estimate θ , we aim to maximize $\mathcal{L}(\theta; y, x_{\text{obs}}) = p(y, x_{\text{obs}}; \theta)$.

A.2. Metropolis–Hastings sampling

During SAEM iterations, Metropolis–Hastings sampling is performed as in Algorithm 1, with target distribution $f(x_{i,\text{mis}}) = p(x_{i,\text{mis}}|x_{i,\text{obs}}, y_i; \theta)$ and proposal distribution $g(x_{i,\text{mis}}) = p(x_{i,\text{mis}}|x_{i,\text{obs}}; \mu, \Sigma)$.

A.3. Calculation of the observed information matrix

Procedure 2 shows how we calculate the observed information matrix.

Algorithm 1 Metropolis-Hastings sampling.

Input: An initial sample $x_{i,\text{mis}}^{(0)} \sim g(x_{i,\text{mis}})$;
for $s = 1, 2, \dots, S$ **do**
 Generate $x_{i,\text{mis}}^{(s)} \sim g(x_{i,\text{mis}})$;
 Generate $u \sim \mathcal{U}[0, 1]$;
 Calculate the ratio $w = \frac{f(x_{i,\text{mis}}^{(s)})/g(x_{i,\text{mis}}^{(s)})}{f(x_{i,\text{mis}}^{(s-1)})/g(x_{i,\text{mis}}^{(s-1)})}$;
 if $u < w$ **then**
 Accept $x_{i,\text{mis}}^{(s)}$;
 else
 $x_{i,\text{mis}}^{(s)} \leftarrow x_{i,\text{mis}}^{(s-1)}$;
 end if
end for
Output: $(x_{i,\text{mis}}^{(s)}, 1 \leq i \leq n, 1 \leq s \leq S)$.

Procedure 2 Calculation of the observed information matrix.

Input: After drawing MH samples $(x_{i,\text{mis}}^{(s)}, 1 \leq i \leq n, 1 \leq s \leq S)$ for unobserved data $(x_{i,\text{mis}}, 1 \leq i \leq n)$, we have imputed observations, noted as $(z_i^{(s)}, 1 \leq i \leq n, 1 \leq s \leq S)$, where $z_{ij}^{(s)} = x_{i,\text{obs}}$, if x_{ij} is observed; else $z_{ij}^{(s)} = x_{i,\text{mis}}^{(s)}$.
for $n = 1, 2, \dots, n$ **do**
 for $s = 1, 2, \dots, S$ **do**
 Calculate the gradient:

$$\nabla f_{is} = \frac{\partial \mathcal{L}(\theta; x_{i,\text{obs}}, x_{i,\text{mis}}^{(s)}, y_i)}{\partial \beta} = z_i^{(s)} \left(y_i - \frac{\exp(\hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j z_{ij}^{(s)})}{1 + \exp(\hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j z_{ij}^{(s)})} \right);$$

 Calculate the Hessian matrix:

$$H_{is} = \frac{\partial^2 \mathcal{L}(\theta; x_{i,\text{obs}}, x_{i,\text{mis}}^{(s)}, y_i)}{\partial \beta \partial \beta^T} = -z_i^{(s)} z_i^{(s)T} \frac{\exp(\hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j z_{ij}^{(s)})}{(1 + \exp(\hat{\beta}_0 + \sum_{j=1}^p \hat{\beta}_j z_{ij}^{(s)}))^2};$$

 $\Delta_i \leftarrow \frac{1}{s} [(s-1)\Delta_i + \nabla f_{is}];$
 $D_i \leftarrow \frac{1}{s} [(s-1)D_i + H_{is}];$
 $G_i \leftarrow \frac{1}{s} [(s-1)G_i + \nabla f_{is} \nabla f_{is}^T];$
 end for
 $\hat{\mathcal{I}}_S(\hat{\beta}) \leftarrow \hat{\mathcal{I}}_S(\hat{\beta}) - (D_i + G_i - \Delta_i \Delta_i^T);$
end for
Output: $\hat{\mathcal{I}}_S(\hat{\beta})$.

A.4. Logistic regression on a simulated complete dataset

Fig. A.12 shows the ROC curve on a simulated complete dataset. The corresponding AUC (for the training set) is 0.8976.

A.5. Simulation results for missing-at-random data

We consider a missing-at-random mechanism to generate data. Fig. A.13 shows that the biases were very similar to those obtained under a MCAR mechanism, and parameters were estimated without bias.

A.6. Simulation results for model misspecification: coverage

Table A.6 shows the coverage for all parameters, and the average lengths of the corresponding confidence intervals in parentheses.

A.7. Definitions of variables in the TraumaBase dataset

In this section, we define the selected quantitative variables:

- Age: Age.
- Poids: Weight.
- Taille: Height.




Fig. A.12. ROC curve on a simulated complete dataset.


Fig. A.13. Empirical distribution of the bias of $\hat{\beta}_3$ obtained under an MAR mechanism, with $n = 1000$ and 10% missing values.

Table A.6

Coverage (%) for $n = 1000$, MCAR, and misspecified models, calculated over 1000 simulations. Bold indicates under-coverage. Inside the parentheses is the average length of the corresponding confidence interval over 1000 simulations (multiplied by 100).

Parameter	no NA	CC	mice	SAEM
Student distribution: ($v = 5$)				
β_0	94.7 (68.02)	94.3 (84.14)	94.6 (67.69)	93.8 (68.25)
β_1	95.2 (54.78)	94.2 (72.15)	91.7 (61.96)	93.5 (63.05)
β_2	94.9 (27.66)	94.6 (36.39)	91.4 (31.21)	93.7 (31.84)
β_3	94.9 (26.76)	94.3 (35.24)	81.5 (30.46)	94.7 (29.98)
β_4	95.2 (11.52)	95.4 (15.16)	95.8 (12.94)	95.5 (12.88)
β_5	93.7 (17.63)	94.9 (23.22)	83.4 (20.40)	93.3 (19.93)
Gaussian mixture:				
β_0	94.8 (57.54)	95.2 (75.42)	95.4 (61.95)	95.0 (61.33)
β_1	94.7 (58.00)	96.2 (76.05)	95.4 (66.66)	95.3 (66.13)
β_2	94.3 (28.49)	95.3 (37.35)	95.3 (32.65)	94.0 (32.50)
β_3	94.7 (26.16)	94.9 (34.38)	94.9 (28.91)	94.5 (29.10)
β_4	94.4 (12.68)	94.4 (16.60)	94.4 (14.24)	94.7 (14.09)
β_5	95.3 (17.70)	94.7 (23.25)	94.7 (19.86)	95.3 (19.92)

- *BMI*: Body Mass index, $BMI = \frac{\text{Weight in kg}}{(\text{Height in m})^2}$
- *Glasgow*: Glasgow Coma Scale.
- *Glasgow.moteur*: Glasgow Coma Scale motor component.
- *PAS.min*: The minimum systolic blood pressure.

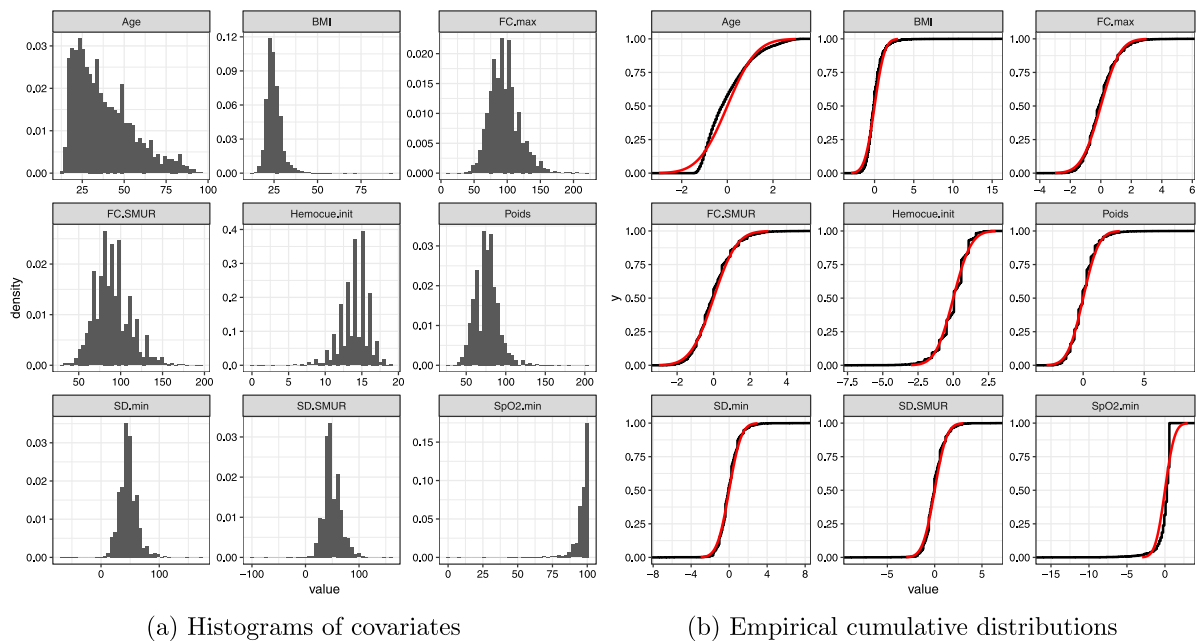


Fig. A.14. Empirical distributions of variables from TraumaBase. (a) Histograms of covariates; (b) The black line is the empirical cumulative distribution while the red one corresponds to the normal distribution.

- *PAD.min*: The minimum diastolic blood pressure.
- *FC.max*: The maximum number of heart beats per unit time (usually a minute).
- *PAS.SMUR*: Systolic blood pressure at ambulance arrival.
- *PAD.SMUR*: Diastolic blood pressure at ambulance arrival.
- *FC.SMUR*: Heart rate at ambulance arrival.
- *Hemocue.init*: Capillary hemoglobin concentration.
- *SpO2.min*: Oxygen saturation.
- *Remplissage.total.colloides* (or *RT.colloides*): Fluid expansion colloids.
- *Remplissage.total.cristalloides* (or *RT.cristalloides*): Fluid expansion cristalloides.
- *SD.min* ($= PAS.min - PAD.min$): Pulse pressure for the minimum values of diastolic and systolic blood pressure.
- *SD.SMUR* ($= PAS.SMUR - PAD.SMUR$): Pulse pressure at ambulance arrival.

Fig. A.14 shows the histogram and the empirical cumulative distribution function of some of the covariates in TraumaBase. Several of these are not symmetric. In practice, it is possible to consider that suitable transformation of covariates can be approximated by normal distributions. For example, transformations of the form $\log(c+x)$ and $\log(c-x)$, may be appropriate for right-skewed and left-skewed distributions respectively. We applied the proposed methodology to the real dataset after the log-transformation. However, the prediction results from cross-validation did not show any advantage to transforming the variables. Indeed when a log transformation is used as a preprocessing step, it only operates on the observed part, which is appropriate under MCAR values. Consequently, we have decided not to use any transformation on the dataset.

A.8. Details on the predictive performance on TraumaBase data

Details on the predictive performance on TraumaBase data are given in Table A.7.

Appendix B. Supplementary material

R-package: The R package *misaem* containing the implementation of the SAEM algorithm to fit the logistic regression model with missing data is available on CRAN (Jiang, 2019a).

Codes: Code to reproduce the experiments is provided on GitHub (Jiang, 2019b).

Additional supplementary materials: Additional simulation results can be found here: (Jiang, 2019c).

Table A.7

Comparisons of the mean of the predictive performance (values are multiplied by 100) of different methods that can deal with missing data. AUC is the area under the ROC curve; the accuracy is the number of true positives plus true negatives, divided by the total number of observations; the sensitivity is the true positive rate; the specificity is the true negative rate; the precision is the number of true positives over all positive predictions. Best results are shown in bold.

Metrics	SAEM	missForest	impMean	impPCA	mice	predRF	predSVM
AUC	88.5	88.8	88.9	89.0	87.7	88.0	80.4
Accuracy	86.9	87.0	87.3	86.7	85.3	87.2	88.3
Precision	41.1	41.6	42.2	41.0	37.9	41.6	44.0
Sensitivity	74.6	74.3	73.2	75.0	75.2	71.5	66.0
Specificity	88.2	88.4	88.8	87.9	86.4	88.9	90.6

References

- Brier, G.W., 1950. Verification of forecasts expressed in terms of probability. *Mon. Weather Rev.* 78 (1), 1–3.
- Chow, W.K., 1979. A Look at Various Estimators in Logistic Models in the Presence of Missing Values. Technical Report, Rand Corp Santa Monica CA.
- Claeskens, G., Consentino, F., 2008. Variable selection with incomplete covariate data. *Biometrics* 64, 1062–1069.
- Consentino, F., Claeskens, G., 2011. Missing covariates in logistic regression, estimation and distribution selection. *Stat. Model.* 11 (2), 159–183.
- Delyon, B., Lavielle, M., Moulines, E., 1999. Convergence of a stochastic approximation version of the EM algorithm. *Ann. Statist.* 27 (1), 94–128.
- Dempster, A.P., Laird, N.M., Rubin, D.B., 1977. Maximum likelihood from incomplete data via the EM algorithm. *J. R. Stat. Soc. Ser. B Stat. Methodol.* 39 (1), 1–38.
- Gilks, W.R., Wild, P.P., 1992. Adaptive rejection sampling for Gibbs sampling. *Appl. Stat.* 41 (2), 337–348.
- Good, I.J., 1952. Rational decisions. *J. R. Stat. Soc. Ser. B Stat. Methodol.* 107–114.
- Hamada, S.R., Gauss, T., Duchateau, F.-X., Truchot, J., Harrois, A., Raux, M., Duranteau, J., Mantz, J., Paugam-Burtz, C., 2014. Evaluation of the performance of french physician-staffed emergency medical service in the triage of major trauma patients. *J. Trauma Acute Care Surg.* 76 (6), 1476–1483.
- Hamada, S.R., Gauss, T., Pann, J., Dünser, M.W., Léone, M., Duranteau, J., 2015. European trauma guideline compliance assessment: The ETRAUSS study. *Crit. Care* 19, 423.
- Hamada, S.R., Rosa, A., Gauss, T., Desclefs, J.-P., Raux, M., Harrois, A., Follin, A., Cook, F., Boutonnet, M., Attias, A., Ausset, S., Dhonneur, G., Langeron, O., Paugam-Burtz, C., Pirracchio, R., Riou, B., de St Maurice, G., Vigué, B., Rouquette, A., Duranteau, J., 2018. Development and validation of a pre-hospital “Red Flag” alert for activation of intra-hospital haemorrhage control response in blunt trauma. *Crit. Care* 22 (1), 113.
- Hay, S.I., et al., 2017. Global, regional, and national disability-adjusted life-years (DALYs) for 333 diseases and injuries and healthy life expectancy (HALE) for 195 countries and territories, 1990–2016: a systematic analysis for the global burden of disease study 2016. *Lancet* 390 (10100), 1260–1344.
- Hentges, A.L., Dunsmore, I.R., 1998. Predictive distributions in binary models with missing data. *Comm. Statist. Simulation Comput.* 27 (3), 735–759.
- Ibrahim, J.G., Chen, M.-H., Lipsitz, S.R., 1999. Monte Carlo EM for missing covariates in parametric regression models. *Biometrics* 55, 591–596.
- Ibrahim, J.G., Chen, M.-H., Lipsitz, S.R., Herring, A.H., 2005. Missing-data methods for generalized linear models: A comparative review. *J. Amer. Statist. Assoc.* 100 (469), 332–346.
- Jiang, W., 2019a. *miSAEM: Logistic regression with missing covariates*. R package version 0.9.1.
- Jiang, W., 2019b. Codes and implementations for “Logistic Regression with Missing Covariates – Parameter Estimation, Model Selection and Prediction within a Joint-Modeling Framework”. https://github.com/wjiang94/miSAEM_logReg.
- Jiang, W., 2019c. Additional supplementary materials for “Logistic Regression with Missing Covariates – Parameter Estimation, Model Selection and Prediction within a Joint-Modeling Framework”. https://github.com/wjiang94/miSAEM_logReg/tree/master/Supplement.
- Jiang, J., Nguyen, T., Rao, J.S., 2015. The E-MS algorithm: Model selection with incomplete data. *J. Amer. Statist. Assoc.* 110 (511), 1136–1147.
- Josse, J., Husson, F., 2016. *missMDA: A package for handling missing values in multivariate data analysis*. *J. Stat. Softw.* 70 (1), 1–31.
- Josse, J., Prost, N., Scornet, E., Varoquaux, G., 2019. On the consistency of supervised learning with missing values. *arXiv e-prints arXiv:1902.06931*.
- Lavielle, M., 2014. *Mixed Effects Models for the Population Approach: Models, Tasks, Methods and Tools*. Chapman and Hall/CRC.
- Little, R.J., Rubin, D.B., 2002. *Statistical Analysis with Missing Data*, second ed. John Wiley & Sons, Inc.
- Liu, Y., Wang, Y., Feng, Y., Wall, M.M., 2016. Variable selection and prediction with incomplete high-dimensional data. *Ann. Appl. Stat.* 10 (1), 418–450.
- Louis, T.A., 1982. Finding the observed information matrix when using the EM algorithm. *J. R. Stat. Soc. Ser. B Stat. Methodol.* 44 (2), 226–233.
- McLachlan, G., Krishnan, T., 2008. *The EM Algorithm and Extensions*, second ed. In: *Wiley Series in Probability and Statistics*, Wiley, Hoboken, NJ.
- Meng, X.-L., Rubin, D.B., 1991. Using EM to obtain asymptotic variance-covariance matrices: The SEM algorithm. *J. Amer. Statist. Assoc.* 86 (416), 899–909.
- R Core Team, 2017. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria.
- Rubin, D.B., 2009. *Multiple Imputation for Nonresponse in Surveys*, Vol. 307. John Wiley & Sons.
- Schafer, J.L., Schenker, N., 2000. Inference with imputed conditional means. *J. Amer. Statist. Assoc.* 95 (449), 144–154.
- Seaman, S., Galati, J., Jackson, D., Carlin, J., 2013. What is meant by “Missing at Random”? *Statist. Sci.* 28 (2), 257–268.
- Stekhoven, D.J., Bühlmann, P., 2012. *MissForest – non-parametric missing value imputation for mixed-type data*. *Bioinformatics* 28 (1), 112–118.
- van Buuren, S., Groothuis-Oudshoorn, K., 2011. *mice: Multivariate imputation by chained equations in R*. *J. Stat. Softw.* 45 (3), 1–67.
- Wei, G.C.G., Tanner, M.A., 1990. A Monte Carlo implementation of the EM algorithm and the poor man’s data augmentation algorithms. *J. Amer. Statist. Assoc.* 85 (411), 699–704.
- Wood, A.M., White, I.R., Royston, P., 2008. How should variable selection be performed with multiply imputed data? *Stat. Med.* 27 (17), 3227–3246.
- Yuen Fung, K., Wrobel, B.A., 1989. The treatment of missing values in logistic regression. *Biom. J.* 31, 35–47.